

Abhishek Pandey

abhishek230102@gmail.com — +91 9022748721 — linkedin.com/in/abhishek-pandey-217a961b2

PROFESSIONAL SUMMARY

Databricks Certified Data Engineer Professional and Reltio Certified Technical Consultant with 2 years of experience across data engineering, backend software development, and applied AI/ML systems. Skilled in designing and building scalable ETL/ELT pipelines, data platforms, and RESTful backend services in cloud-native environments, with hands-on expertise in Databricks, PySpark, Delta Lake, Apache Kafka, FastAPI, PostgreSQL, and Azure Data Services. Experienced in metadata-driven architectures, event-driven systems, and automated CI/CD workflows, with a track record of building production-grade software independently end-to-end. Growing expertise in AI engineering, including LLM integration, LangChain, retrieval-augmented generation (RAG), and vector search, gained through independent projects. Comfortable operating across data engineering, software engineering, and AI engineering roles.

TECHNICAL SKILLS

Programming: Python, SQL, Data Structures & Algorithms

Software Engineering: System Design, REST API Design, Microservices Architecture, Object-Oriented Design

Version Control: Git, GitHub

Data Engineering: Databricks, PySpark, Spark SQL, Delta Lake, Data Lakehouse, Snowflake, Hadoop, ETL, ELT, Data Modeling, Data Pipelines, Data Warehousing

Backend Development: FastAPI, REST APIs

Databases: PostgreSQL

Streaming: Apache Kafka, Event-Driven Architecture

AI/ML & LLM Engineering: LLM Integration (Grok, OpenAI APIs), LangChain, Retrieval-Augmented Generation (RAG), Vector Databases (Qdrant), Semantic Search & Embeddings (Sentence-Transformers), Quantitative/Technical Analysis (pandas_ta)

Workflow Orchestration: Apache Airflow, APScheduler

Automation & Web Scraping: Apify, GitHub Actions

DevOps & CI/CD: GitHub Actions, Databricks Asset Bundles, Automated Deployments

Master Data Management: Reltio

Cloud & Data Services: Azure Data Factory, ADLS Gen2, Azure SQL

Monitoring & Logging: Splunk

Methodologies: Agile, Scrum

PROFESSIONAL EXPERIENCE

Modak Analytics (Client: Humana)

India

Software Development Engineer

Oct 2024 – Present

Project: UDAP WILD – Enterprise Data Product Platform

- Led backend development of UDAP WILD, an enterprise self-service Data Product Platform enabling governed work intake, metadata management, approval workflows, and lifecycle management for organization-wide data products, as part of a 4-member cross-functional team.
- Delivered 4 core platform capabilities end-to-end: File Upload, Dataset Management (Databricks catalog, schema, and table discovery), Data Archival, and Bulk Databricks Job Management (resume, pause, and cancel job runs).
- Designed and developed scalable RESTful APIs using Python and FastAPI with PostgreSQL as a centralized metadata repository, supporting file uploads up to 20GB and validated for 100+ concurrent uploads across 2,000+ daily platform users.
- Built and optimized metadata synchronization ETL pipelines using Databricks and PySpark to process catalog, schema, and table-level metadata across environments.
- Architected Kafka-based event triggers connecting UDAP WILD to ADLS, automatically initiating Data Archival workflows on receipt of designated event messages, improving system decoupling and scalability.
- Designed data archival and lifecycle management workflows with policy-based automation on Databricks tables to ensure governance and compliance.

- Implemented CI/CD automation for Databricks job deployments using Asset Bundles and GitHub Actions, eliminating manual deployment steps and ensuring version-controlled releases.
- Integrated centralized logging by streaming Databricks job logs into Splunk dashboards, enhancing monitoring, observability, and incident resolution time.

Data Engineering Intern

June 2024 – Oct 2024

- Developed and maintained PySpark-based data ingestion and transformation ETL pipelines in Databricks for structured and semi-structured datasets.
- Assisted in metadata catalog management, schema synchronization, and performance optimization of data workflows.
- Contributed to backend API development, unit testing, debugging, and performance tuning to improve service response times.

PROJECTS

AI-Powered Job Newsletter Agent

Personal Project

- Built an end-to-end automated pipeline that scrapes LinkedIn job postings via Apify, ranks them by relevance using semantic similarity search (Sentence-Transformers), and scores/filters candidates using a Groq-hosted LLM (Llama-3.3-70B-Versatile).
- Automated daily execution and delivery through GitHub Actions, generating and emailing an HTML digest of ranked job matches.
- Iterated on filtering logic to improve relevance for the Indian job market, including removing seniority filters that mislabeled listings and adding negative-keyword blocking for irrelevant titles.
- Added a rejected-jobs feedback section to the newsletter to improve transparency and support iterative tuning of the ranking pipeline.

HANDS-ON TRAINING & LABS

Snowflake Data Engineering Labs

Self-Paced Training

- Designed and implemented relational database schemas, tables, views, sequences, and file formats using Snowflake SQL.
- Built data ingestion workflows using COPY INTO, external stages, and custom file formats to load structured and semi-structured datasets.
- Worked with VARIANT data types and JSON parsing to process and transform semi-structured data.
- Wrote SQL using joins, aggregations, window functions, and CTEs to support reporting and analysis use cases.

EDUCATION

Savitribai Phule Pune University

2024

Master of Science (M.Sc.) in Computer Science

CERTIFICATIONS

Databricks Certified Data Engineer Professional

Reltio Certified Technical Consultant